Build a Container Image from Scratch

danishpraka.sh

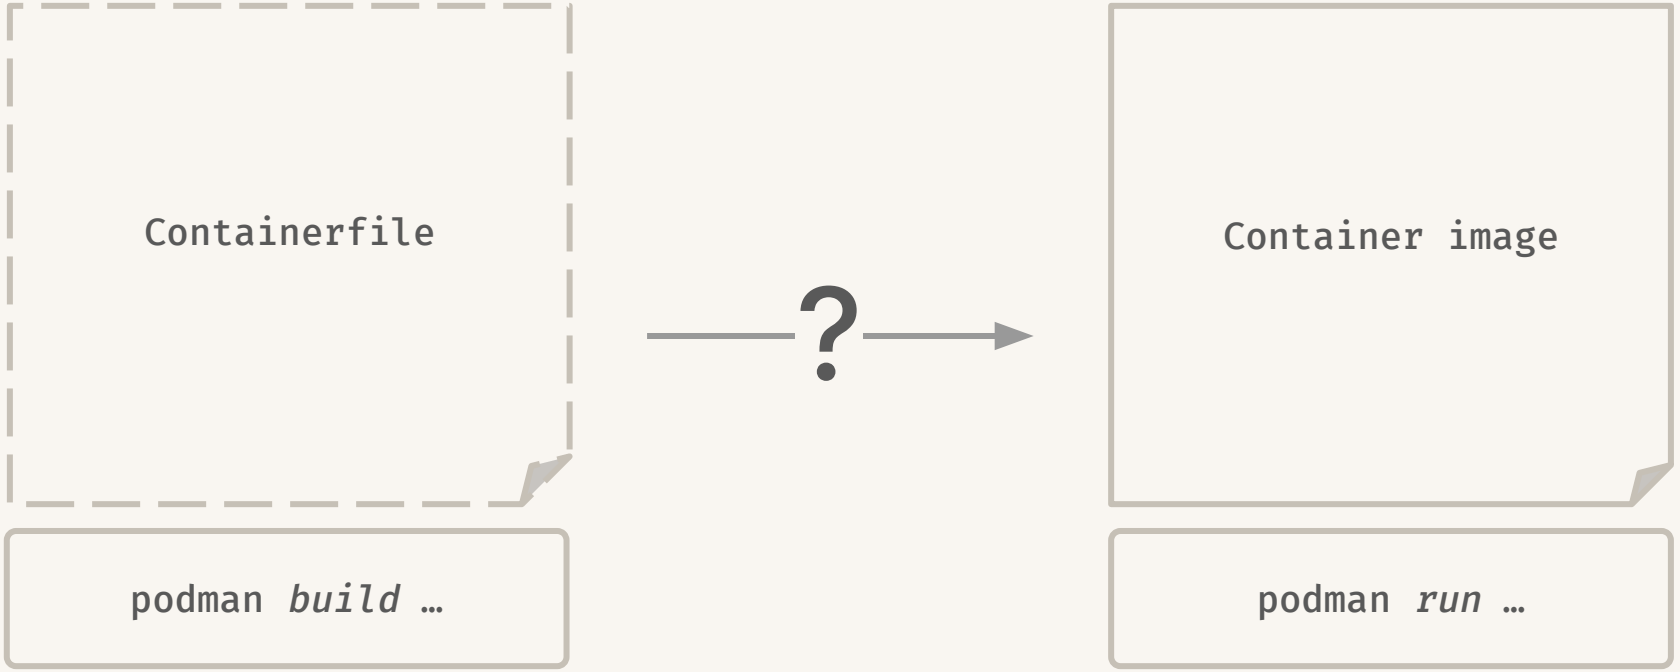
Containerfile
?
Container image
podman *build* …
podman *run* …

container image?

what makes up a container image?

let’s build an image

let’s build an image based on our learnings

demo & verification

demo, load and inspect the image via podman

container image?

OCI image

what constitutes an OCI image

→ Open Containers Initiative

→ Image Specification

→ four major components

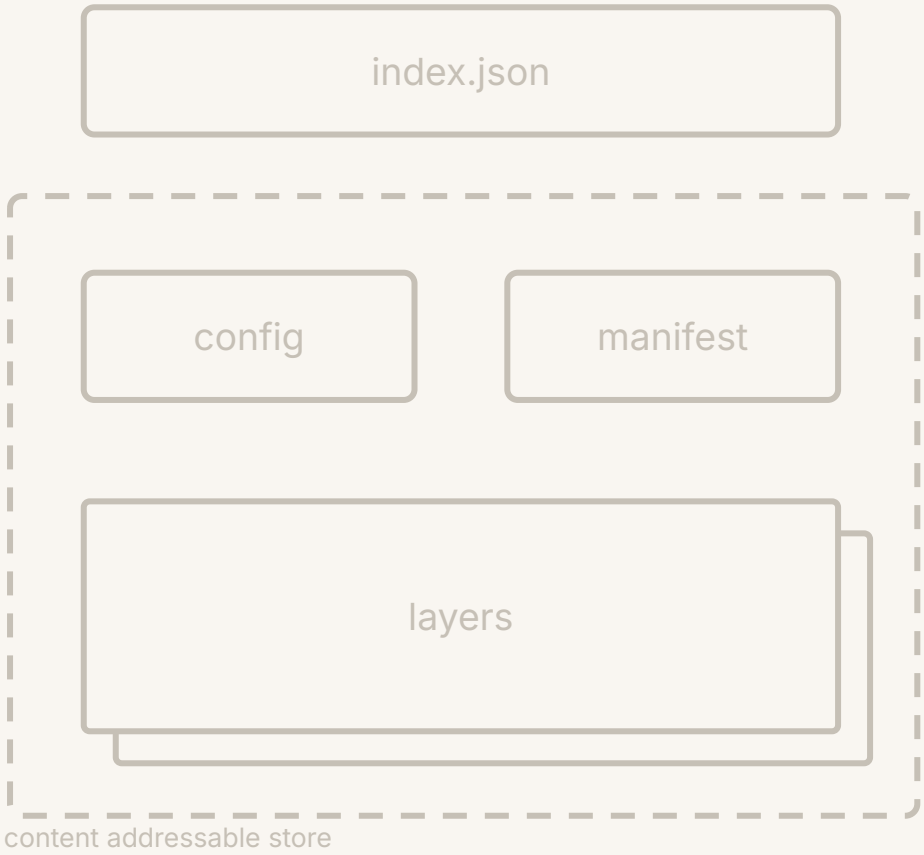

1. layer

what's *inside* the container image

building block for containers
- → filesystem
- → source code

filesystem changeset
- → diff rootfs1..rootfs2

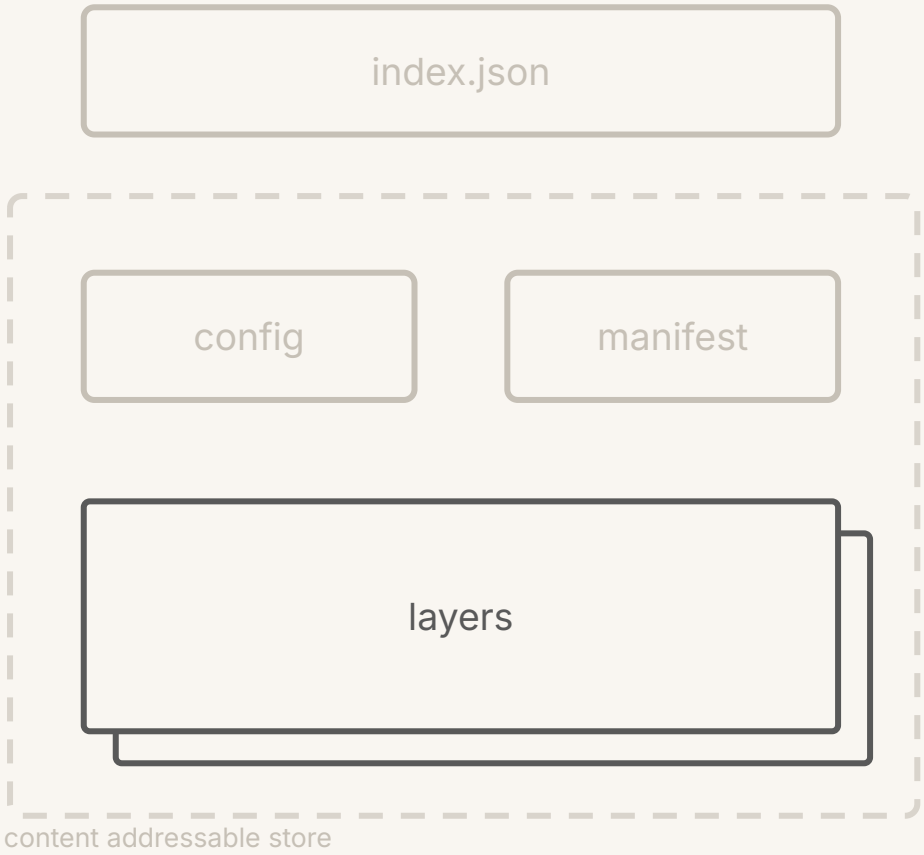

1. layer: create a changeset

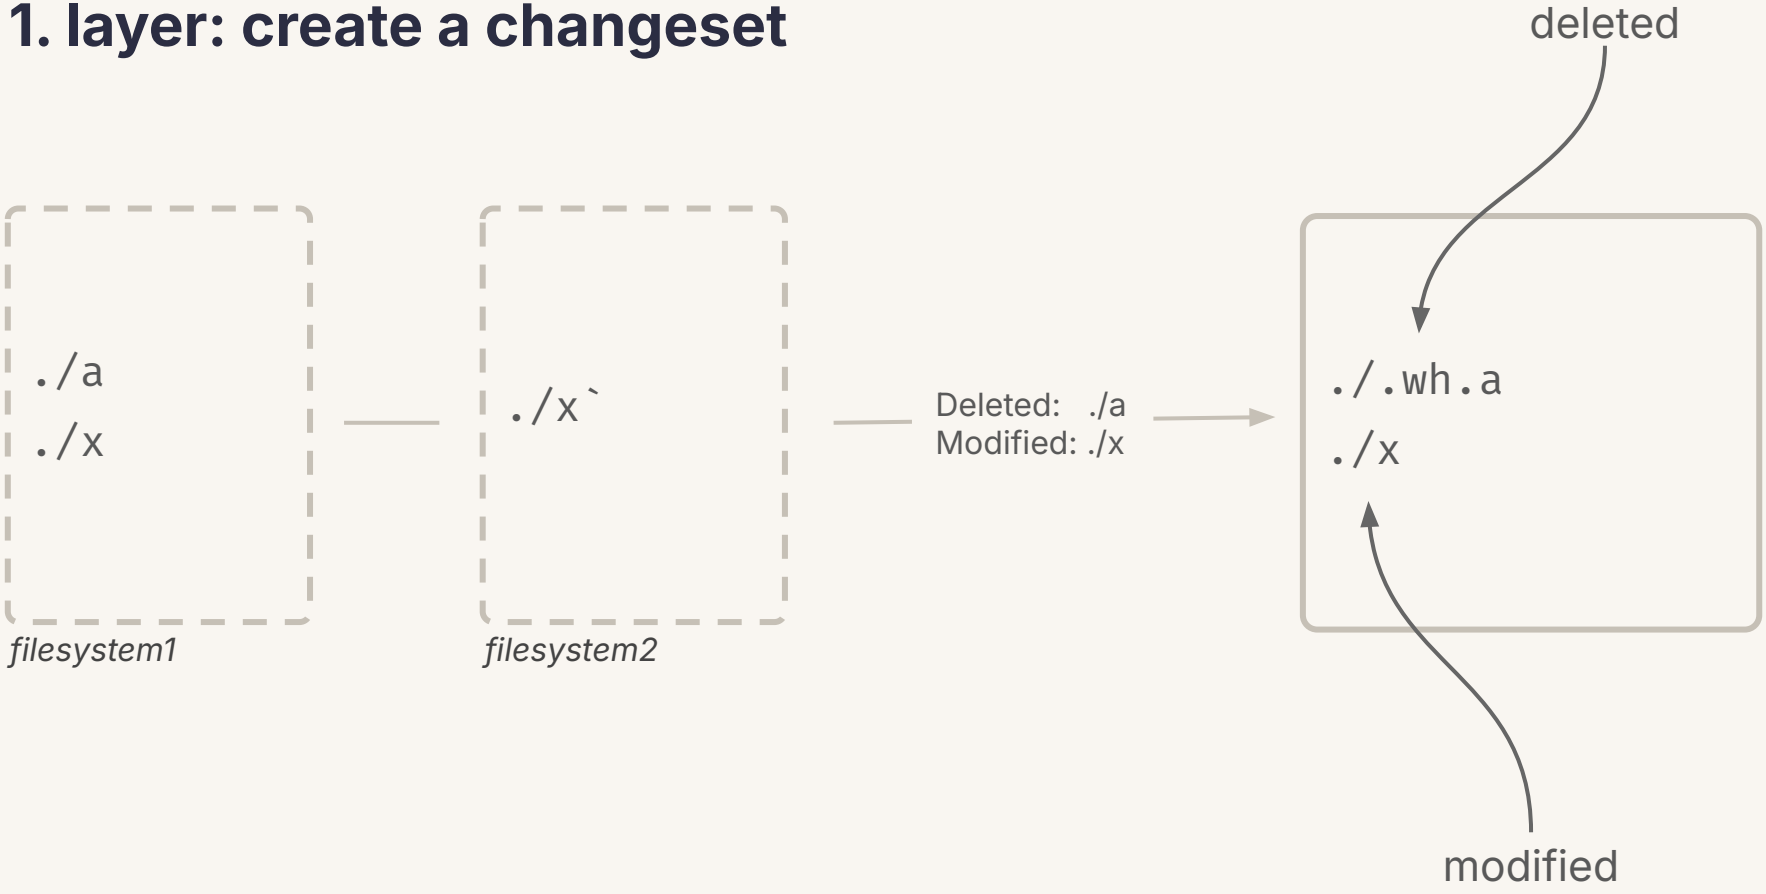

1. layer: create the filesystem

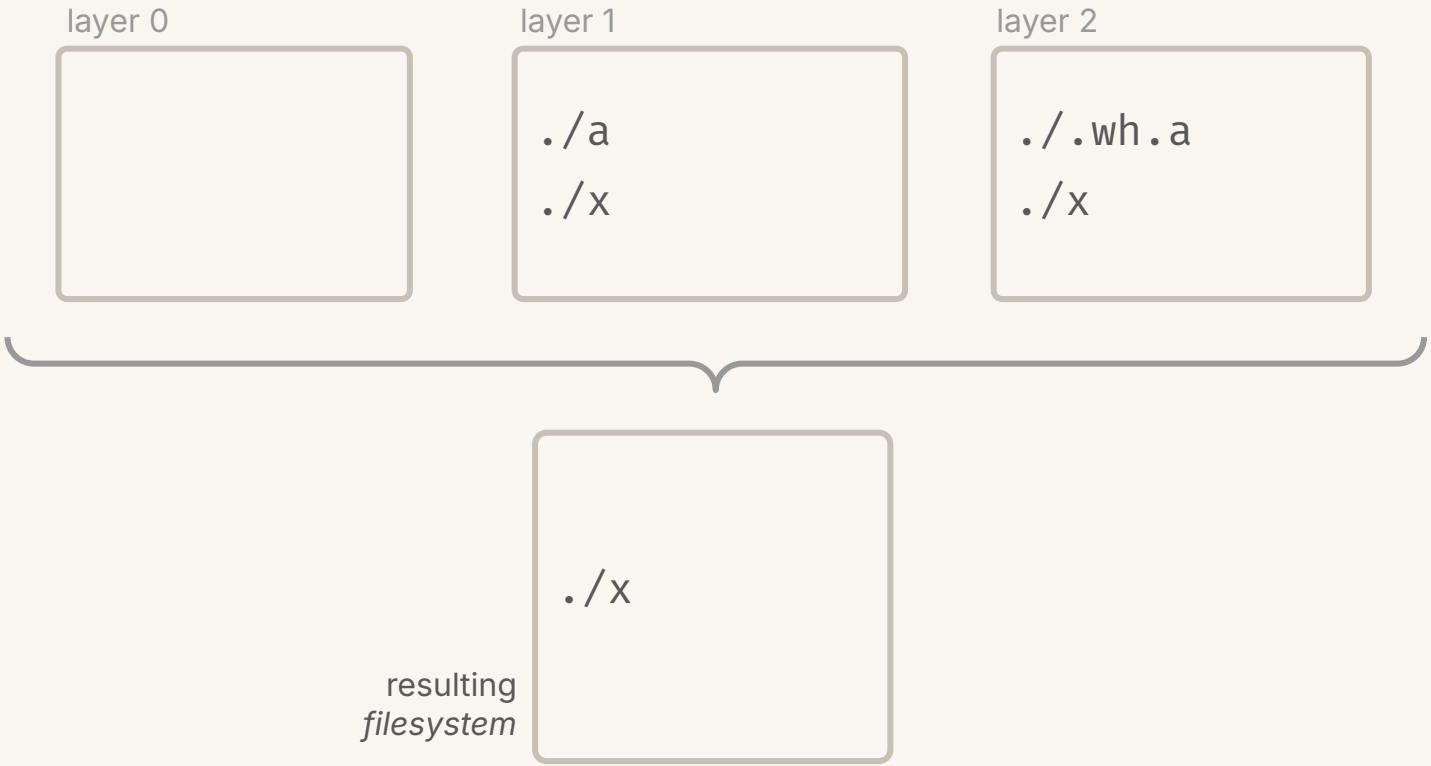

2. config

how to *run* the container

environment variables?

→ HOSTNAME=globalhost

→ PATH=/bin

container entrypoint

→ ./start.sh

container cli options

→ podman run \

--volume

--env

--port

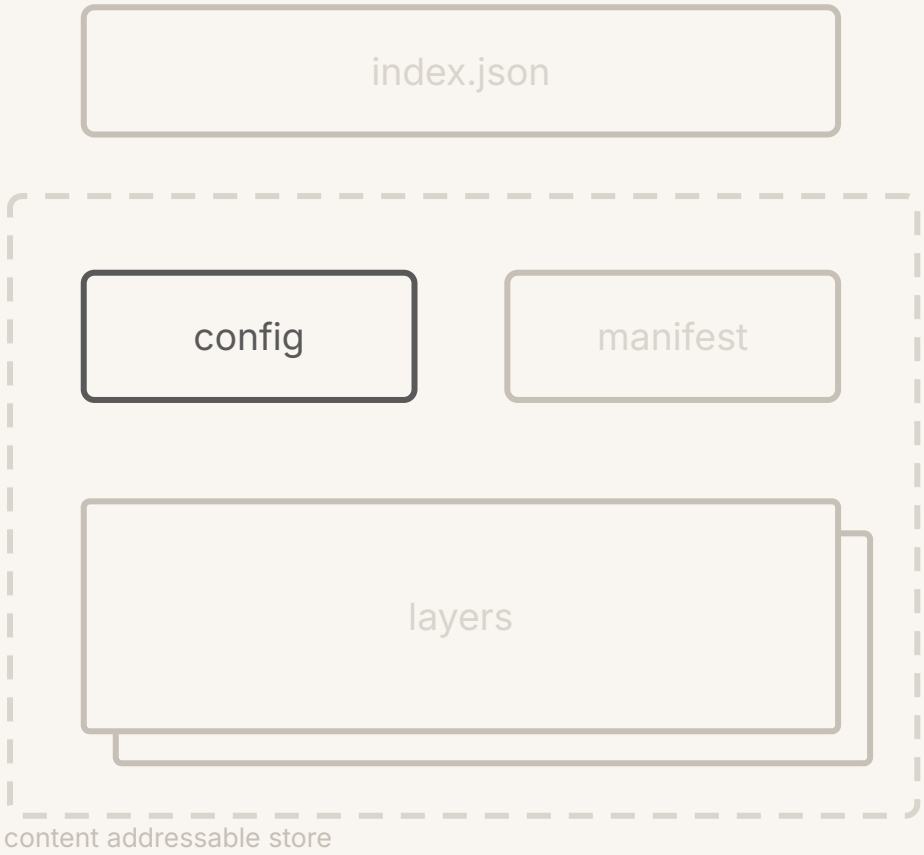

3. manifest

provide config *and* layer info

configuration digest
→ sha256:75b148a9a5b

layer digests
→ sha256:c37c06cdec9

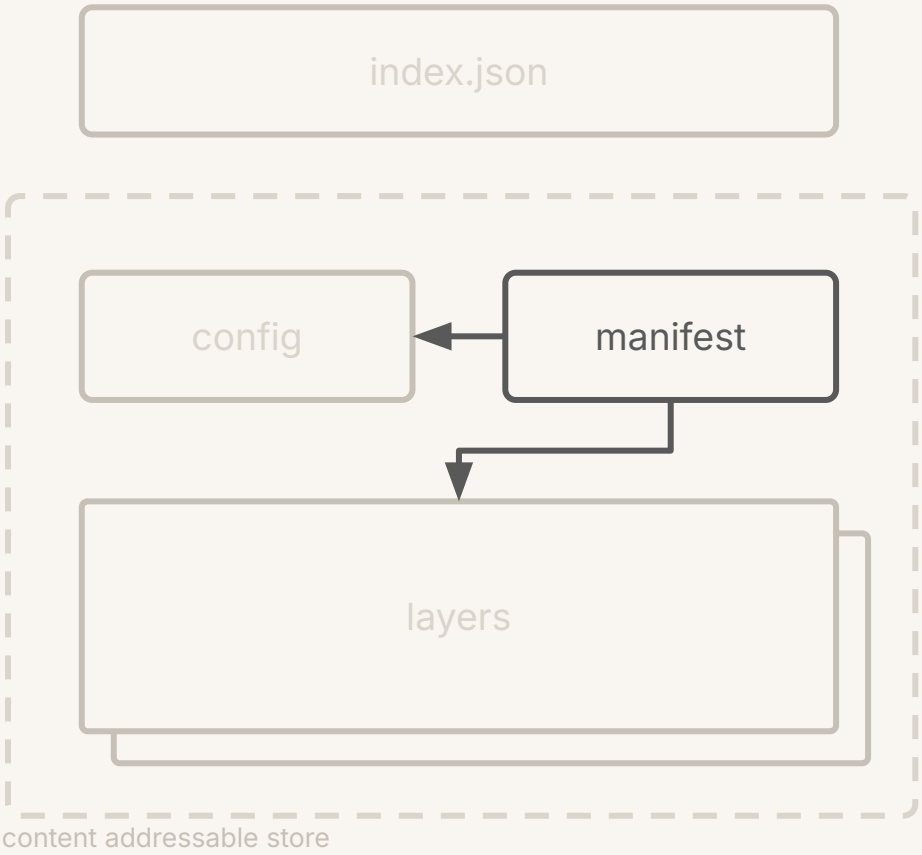

content addressable store

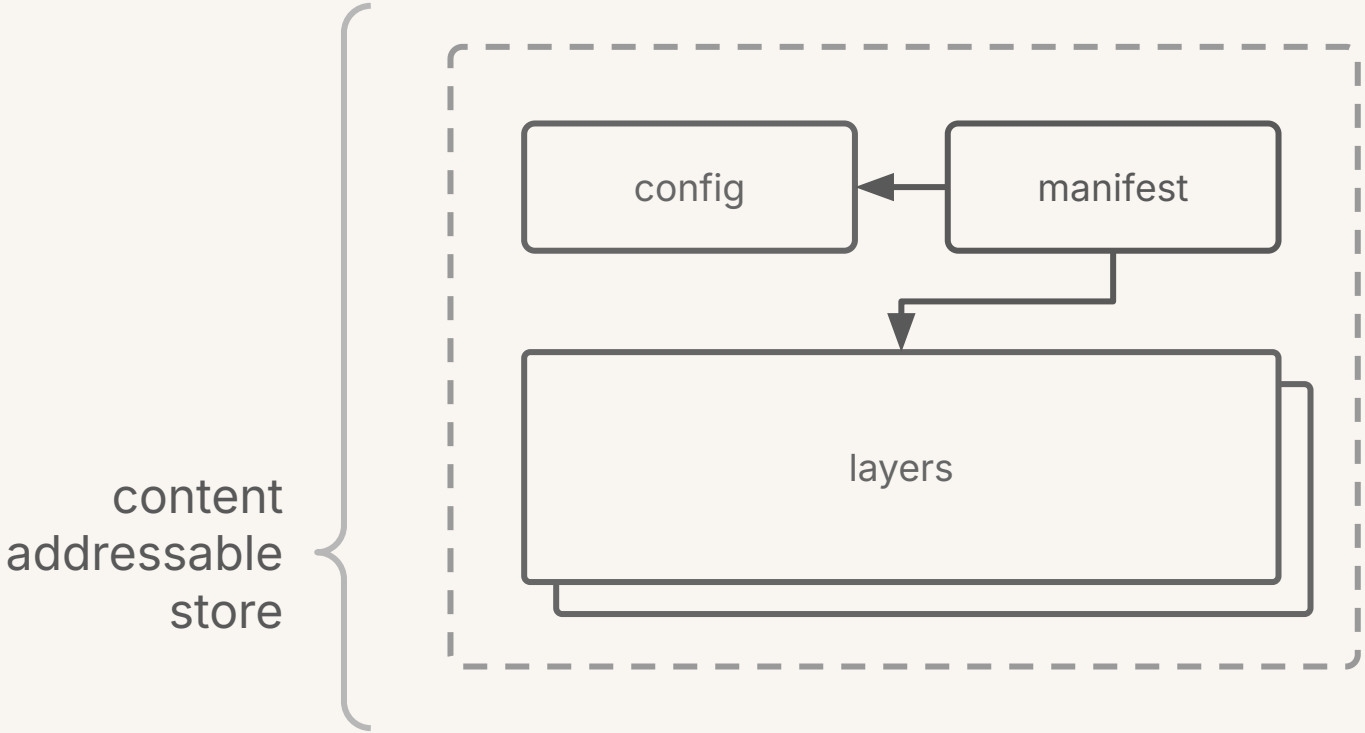

content addressability

```
$ cat event.md
Open Source Summit

$ sha256sum event.md
4c51d469b15d3423dd2d2…

$ mv event.md 4c51d469b15d3423dd2d2…
$ cat 4c51d469b15d3423dd2d2…
Open Source Summit
```

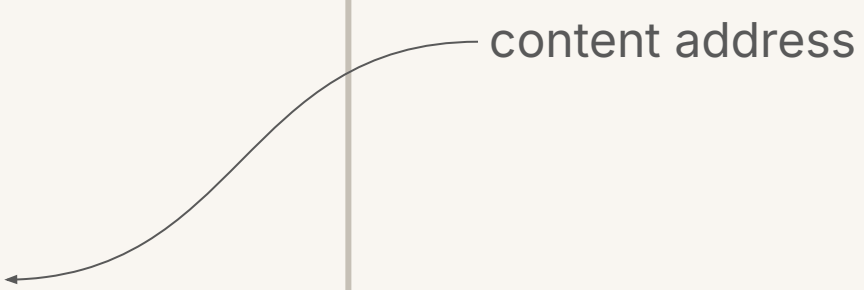

content addressable store

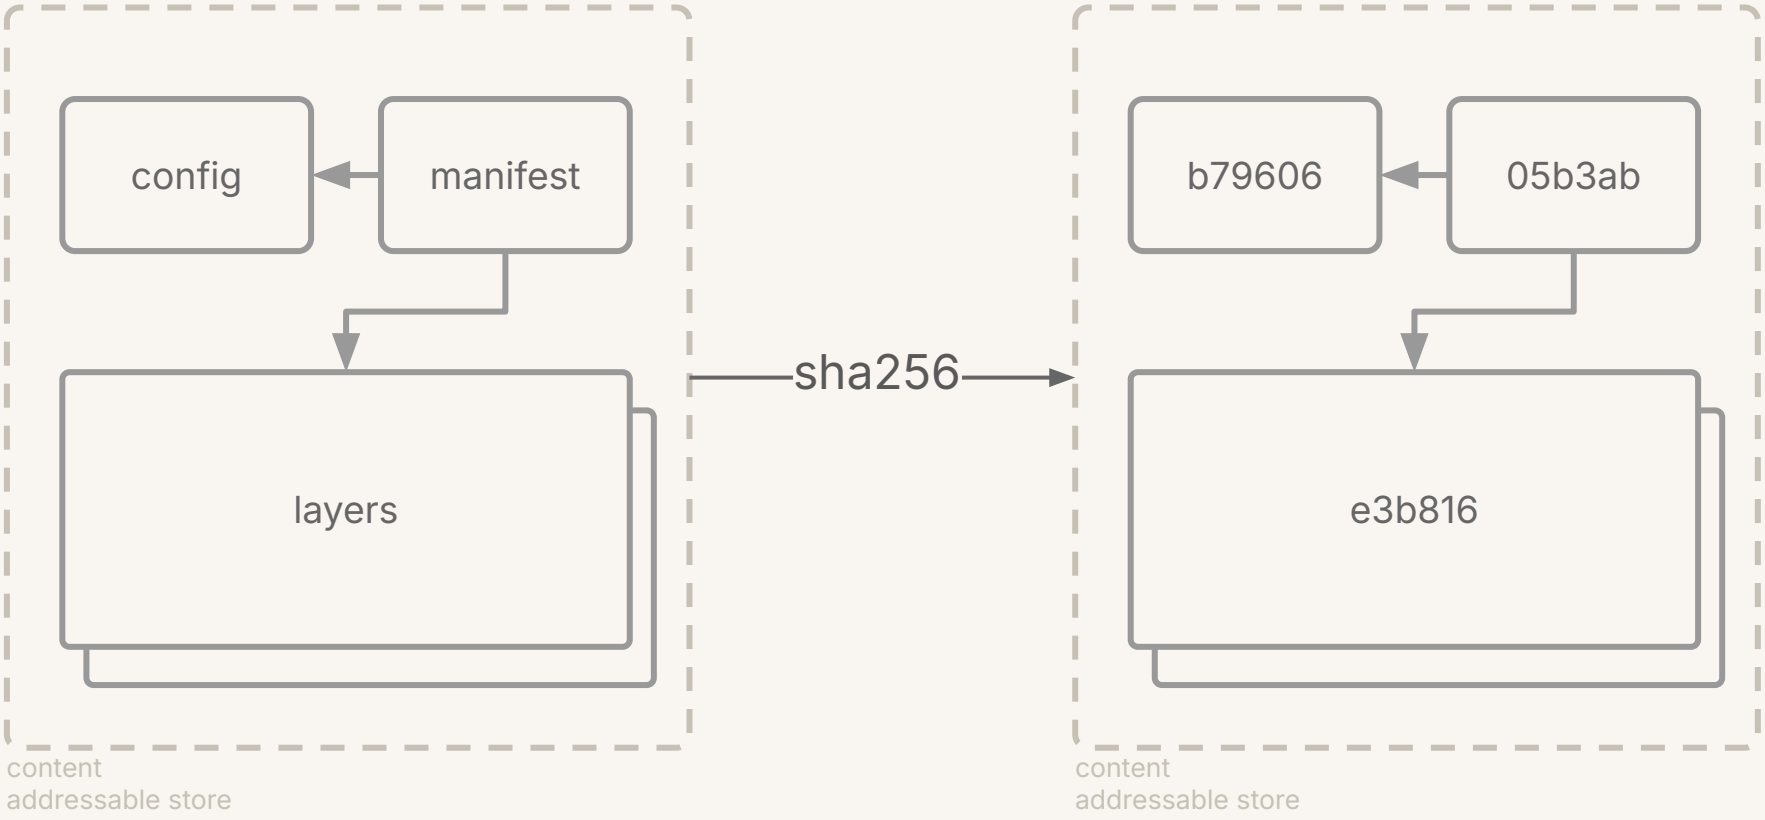

4. index

manifest *for* manifest.json(s)

points to manifest.json

→ sha256:8289bd1bdc2

spans different platforms

→ linux/amd64

→ linux/arm

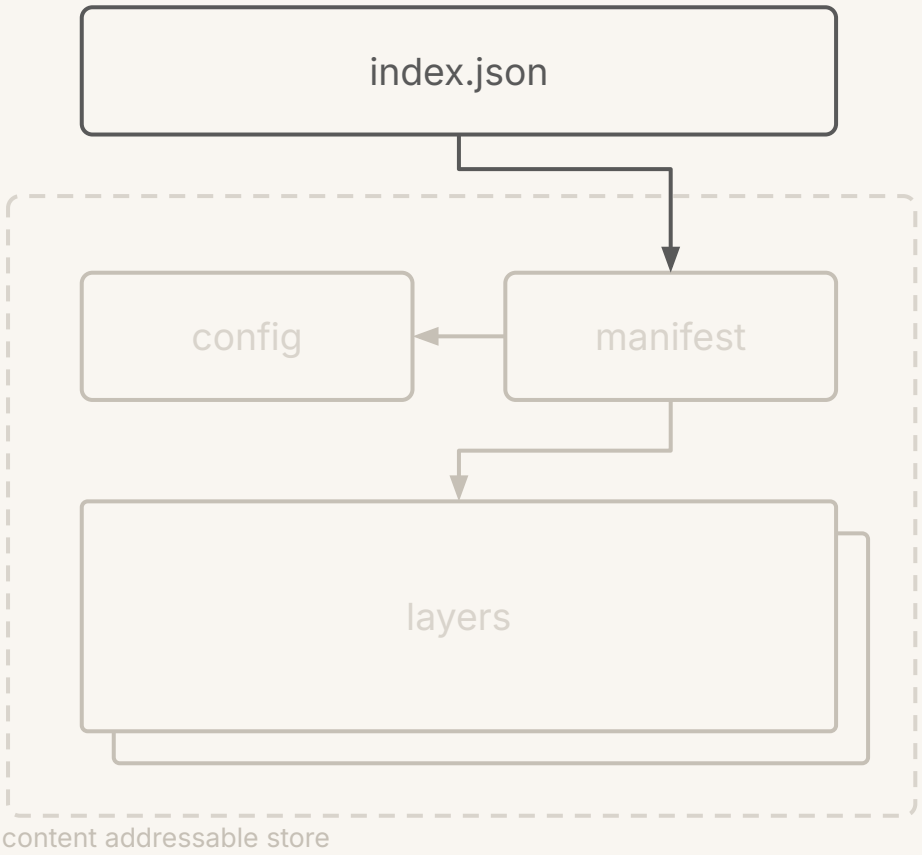

multi-platform

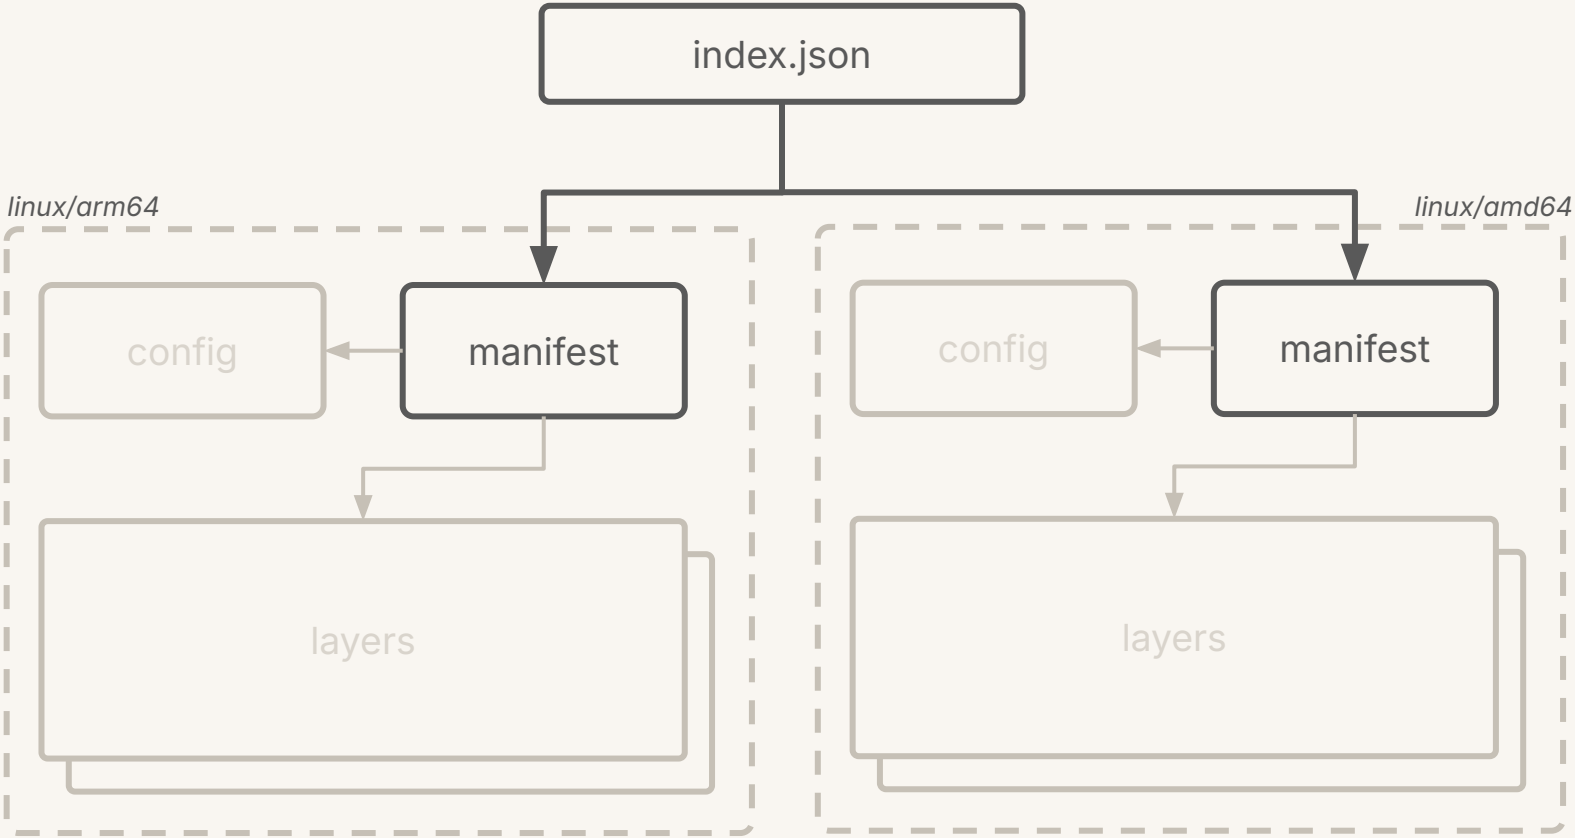

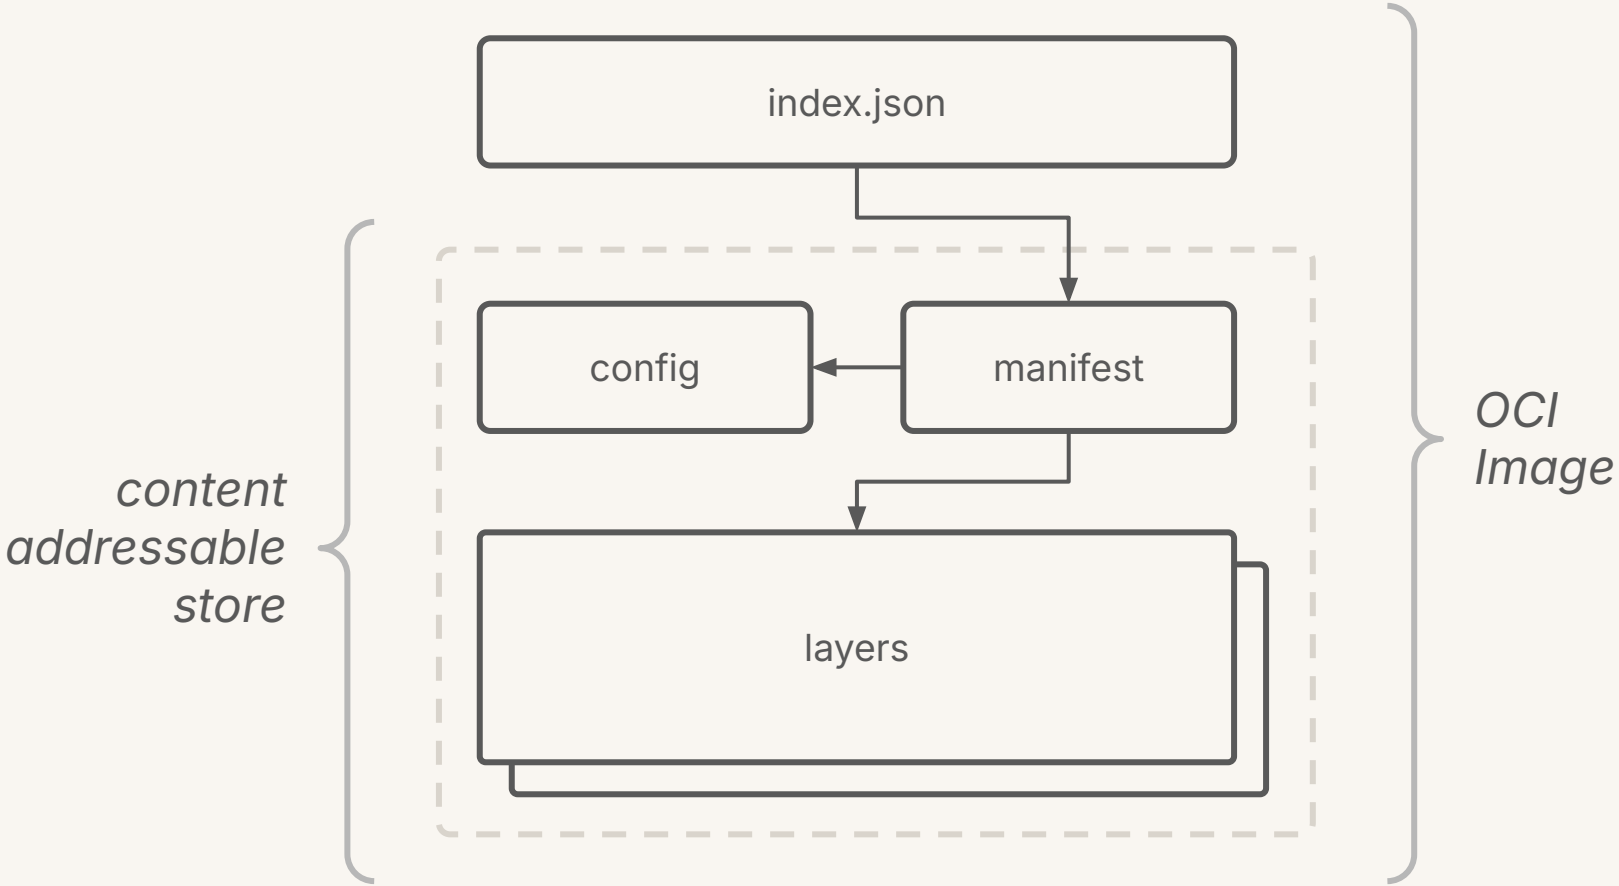
index.json
config
manifest
layers
content addressable store
OCI Image

let’s build an image!

scratch image

first variant with *no* filesystem

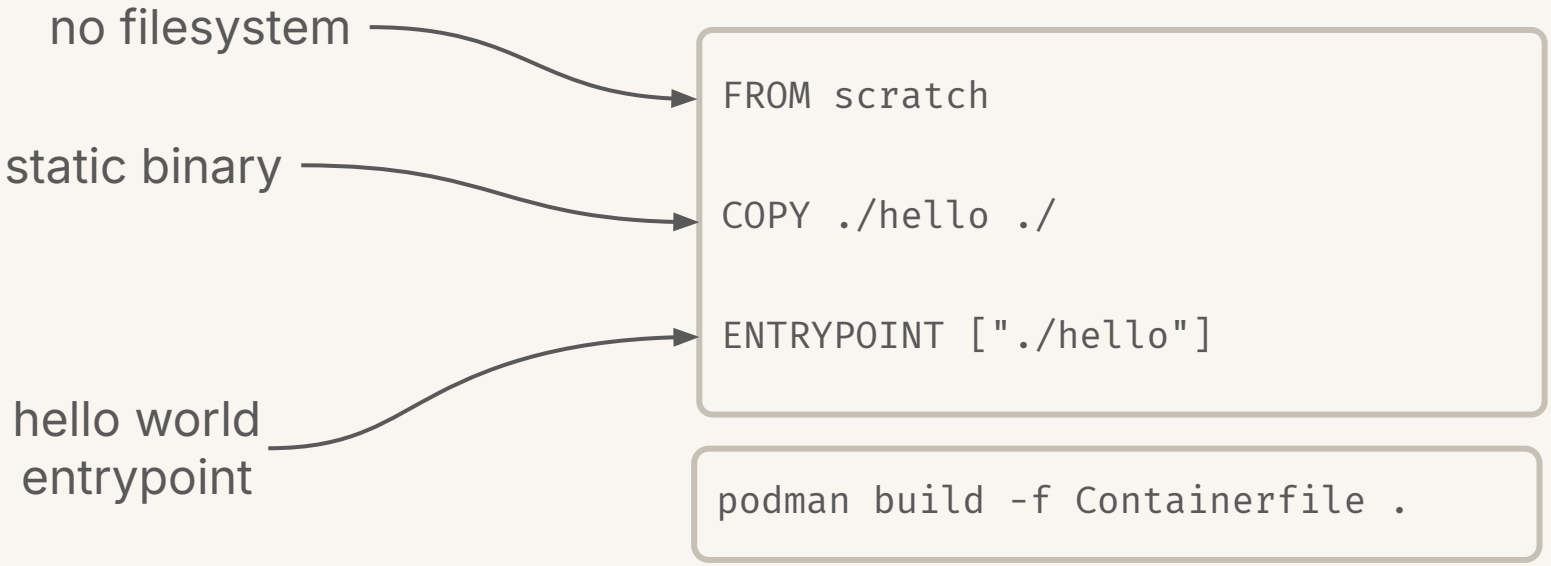

image layout

the OCI image spec layout

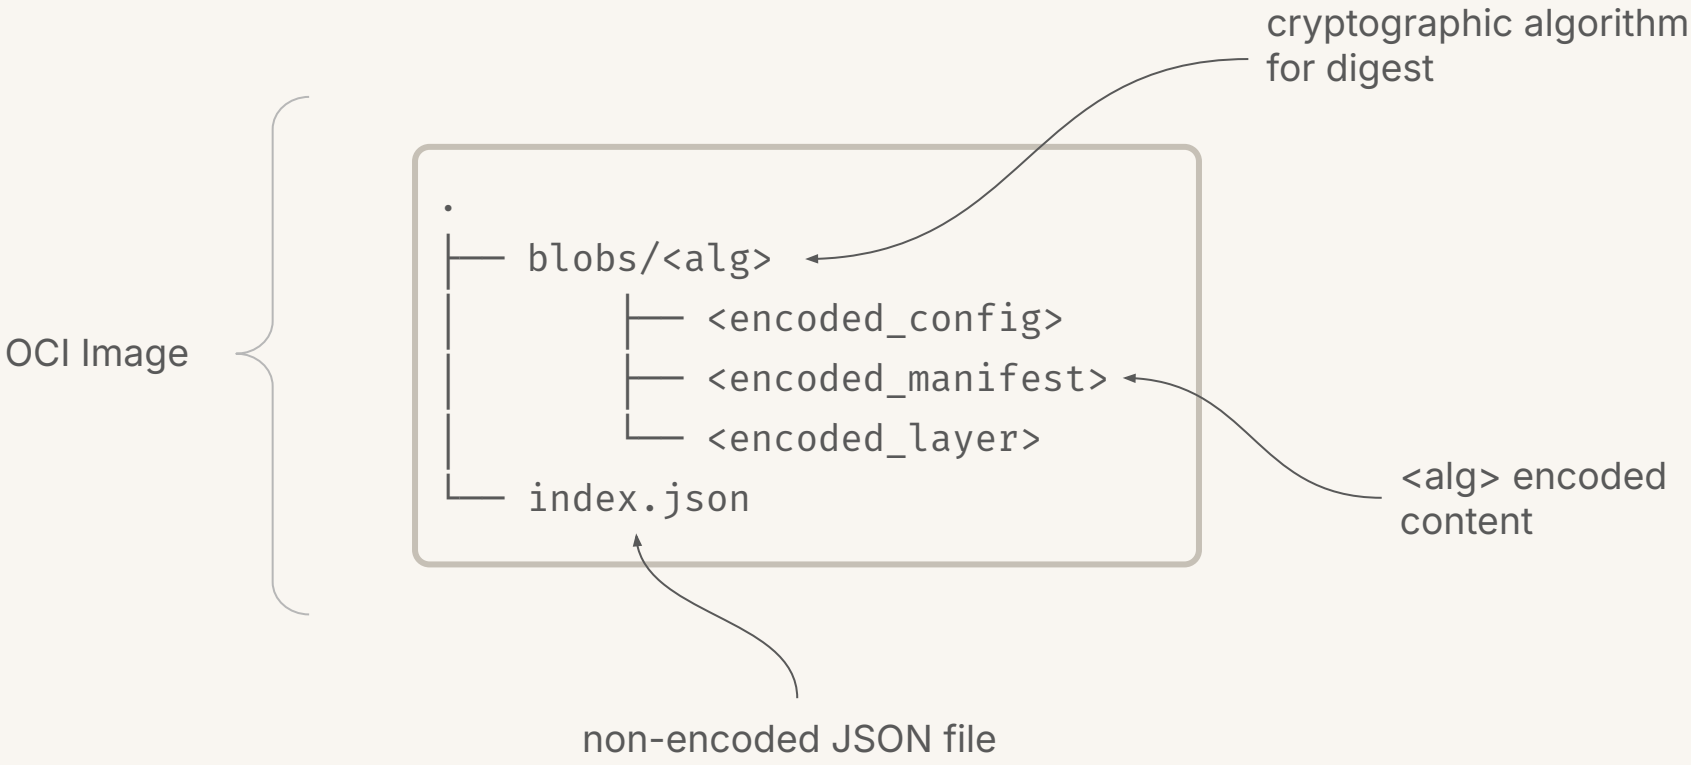

scratch image: layer

first and only layer with a binary

gzipped
tar archive

```
$ gcc -o hello hello.c -static
$ tar -czvf layer.tar.gz hello

$ sha256sum layer.tar.gz
c37c06cdec9d6a0f2a2d5 layer.tar
$ mv layer.tar c37c06cdec9d6a0f2a2d5


$ tree ../..
└── blobs/sha256
    └── c37c06cdec9d6a0f2a2d5
```

content
addressable
archive

layout

scratch image: config

setting the entrypoint and other config options

platform

entrypoint

```
$ cat config.json
{
    "architecture": "amd64",
    "os": "linux",
    "config": {
    "Entrypoint": [
        "./hello"
        ]
    }
}

$ mv config.json $(sha256sum
config.json | awk '{print $1}')
```

scratch image: manifest

identifying the layers and the config

```
$ cat manifest.json
{
  "schemaVersion": 2,
  "mediaType": "application/vnd.oci.image.manifest.v1+json",
  "config": {
     "mediaType": "application/vnd.oci.image.config.v1+json",
     "digest": "sha256:bf7031d43f7d2c9ec77ed",
     "size": 296
  },
  "layers": [
     {
       "mediaType": "application/vnd.oci.image.layer.v1.tar+gzip",
       "digest": "sha256:583eb9106f1be6dffa3f0",
       "size": 1377804
     }
  ]
}

$ mv manifest.json $(sha256sum manifest.json | awk '{print $1}')
```

content descriptor

config

layer digest

scratch image: index

identifying the only manifest in our image

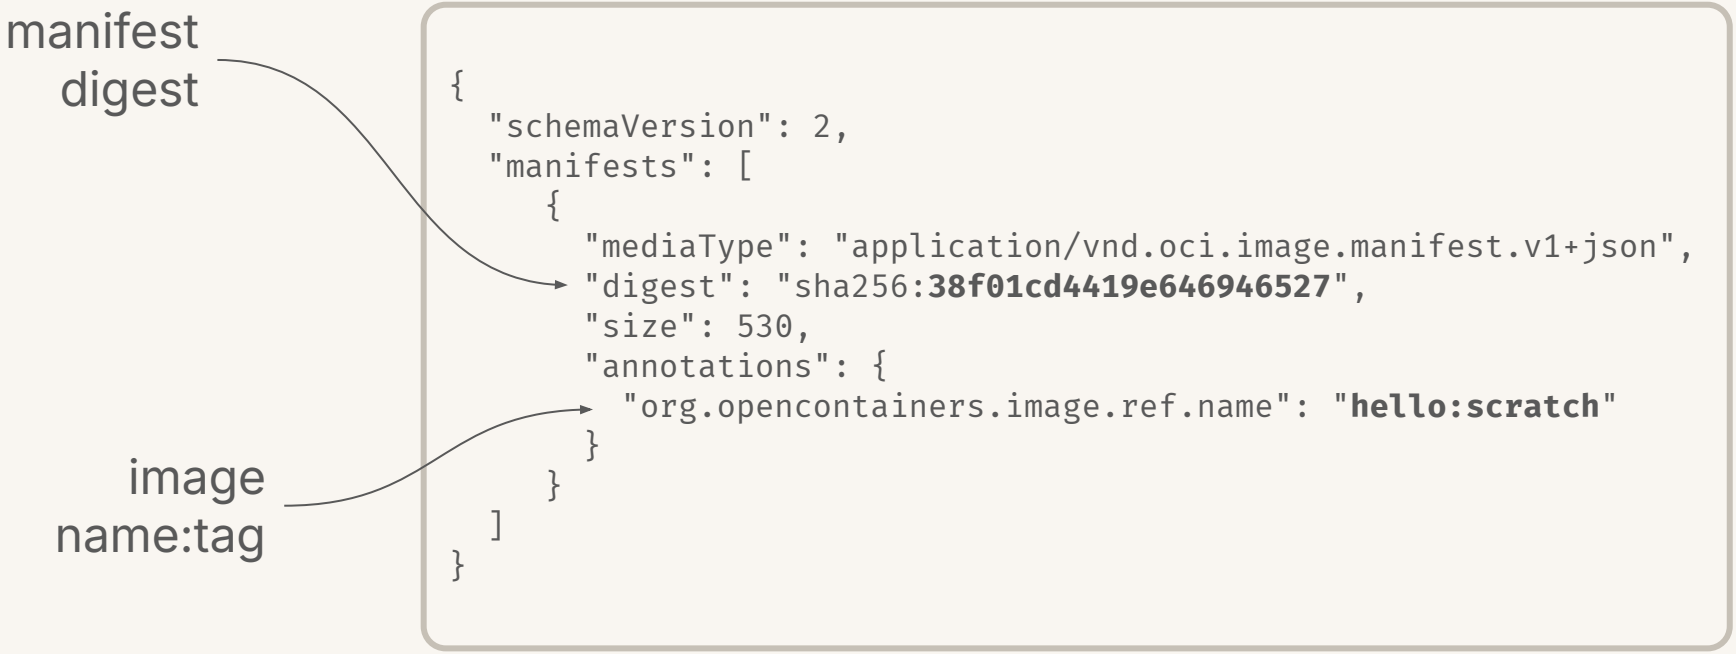

scratch image: load & verify

```
$ tree
├── blobs/sha256
│   ├── 38f01cd4419e646946527
│   ├── c37c06cdec9d6a0f2a2d5
│   └── cd12bca58eb0ed29b24e4
└── index.json

$ tar -cf hello.tar *
```

```
$ podman load < hello.tar
Getting image source signatures
Copying blob c37c06cdec9d done   |
Copying config cd12bca58e done   |
Writing manifest to image destination
Loaded image: localhost/hello:scratch

$ podman image ls hello
REPOSITORY   TAG       IMAGE ID        SIZE
hello        scratch       25e8b3bd9720    3.67MB

$ podman run localhost/hello:scratch world
hello world!
```

demo!

alpine base image

first variant but *with* alpine base image

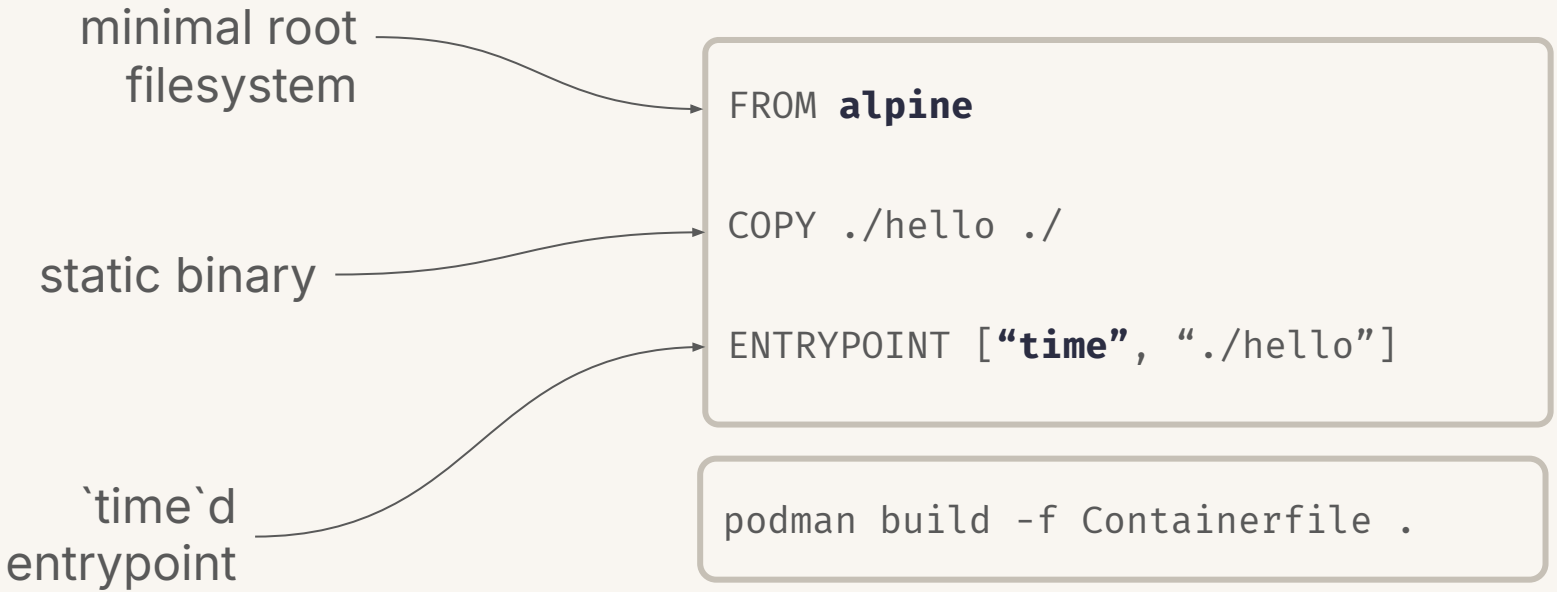

alpine base: layer

using an alpine base root filesystem

```
$ wget https://.../alpine-minirootfs.tar.gz
$ sha256sum alpine-minirootfs.tar.gz
c59d5203bc6b8b6ef81f3



$ tree ../..
└── blobs/sha256
    ├── c37c06cdec9d6a0f2a2d5
    ├── c59d5203bc6b8b6ef81f3
    └── …
```

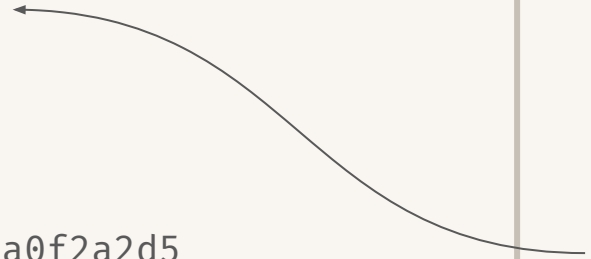

content addressable archive

alpine base: manifest

identifying the layers and the config

```
-snip-

  "layers": [
     {
       "mediaType": "application/vnd.oci.image.layer.v1.tar+gzip",
       "digest": "sha256:c59d5203bc6b8b6ef81f3",
       "size": 3279768
     },
     {
       "mediaType": "application/vnd.oci.image.layer.v1.tar+gzip",
       "digest": "sha256:c37c06cdec9d6a0f2a2d5",
       "size": 1372611
     }
  ]
}
```

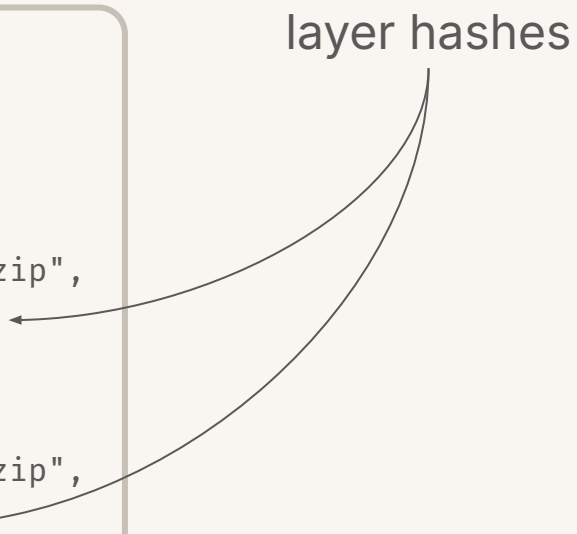

fin

→ container image *internals*

→ how *layers* are used

→ *multi-platform* images

resources

https://github.com/opencontainers/image-spec

https://danishpraka.sh/posts/build-a-container-image-from-scratch/

https://danishpraka.sh/static/build-a-container-image-from-scratch.pdf

danishpraka.sh

Software Engineer
SUSE