

Build a Container Image from Scratch

danishpraka.sh

Containerfile

`podman build ...`



Container image

`podman run ...`

container image?

what makes up a container image?

let's build an image

let's build an image based on our learnings

demo & verification

demo, load and inspect the image via podman

container image?

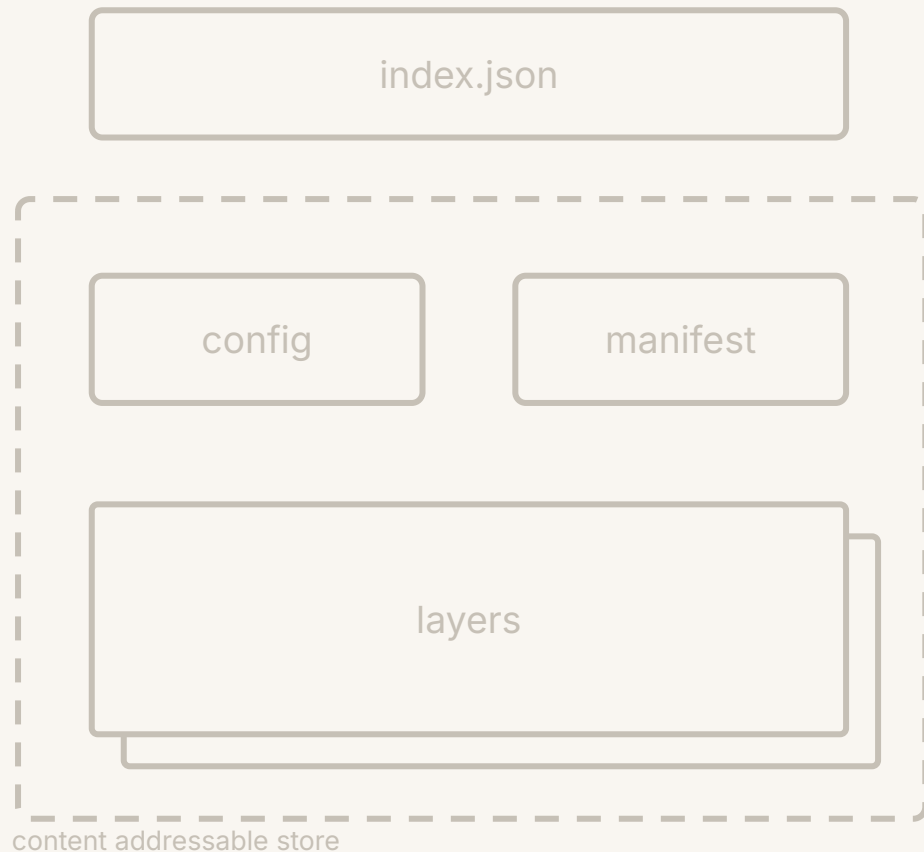
OCI image

what constitutes an OCI image

→ Open Containers Initiative

→ Image Specification

→ four major components



1. layer

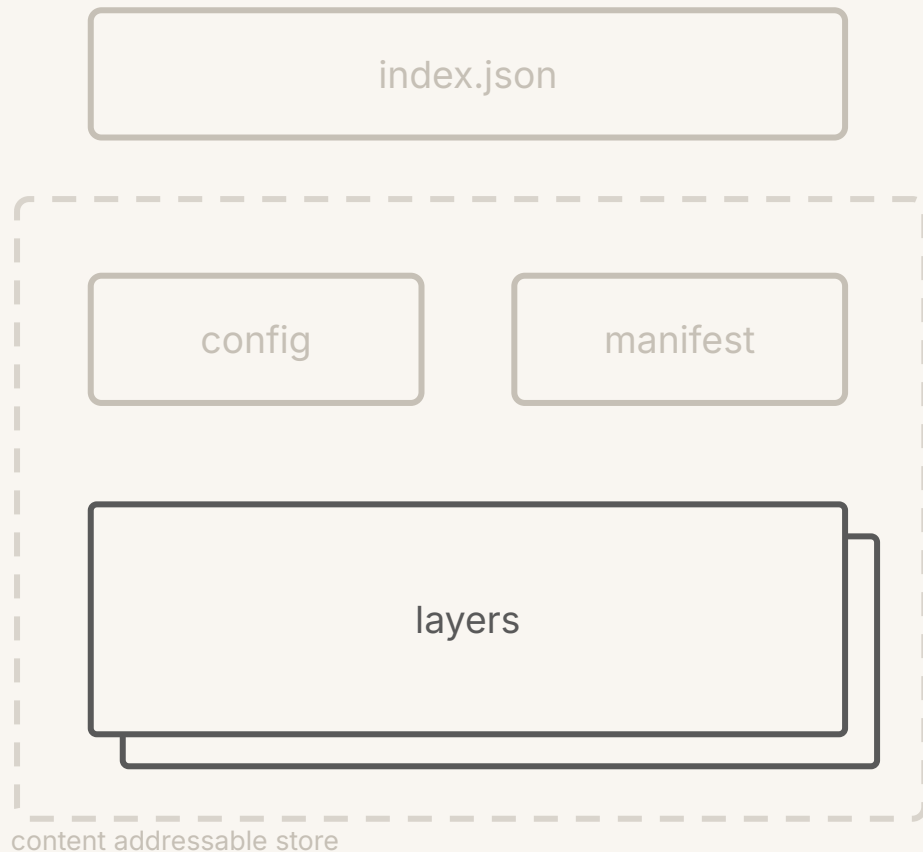
what's *inside* the container image

building block for containers

- filesystem
- source code

filesystem changeset

- diff rootfs1..rootfs2



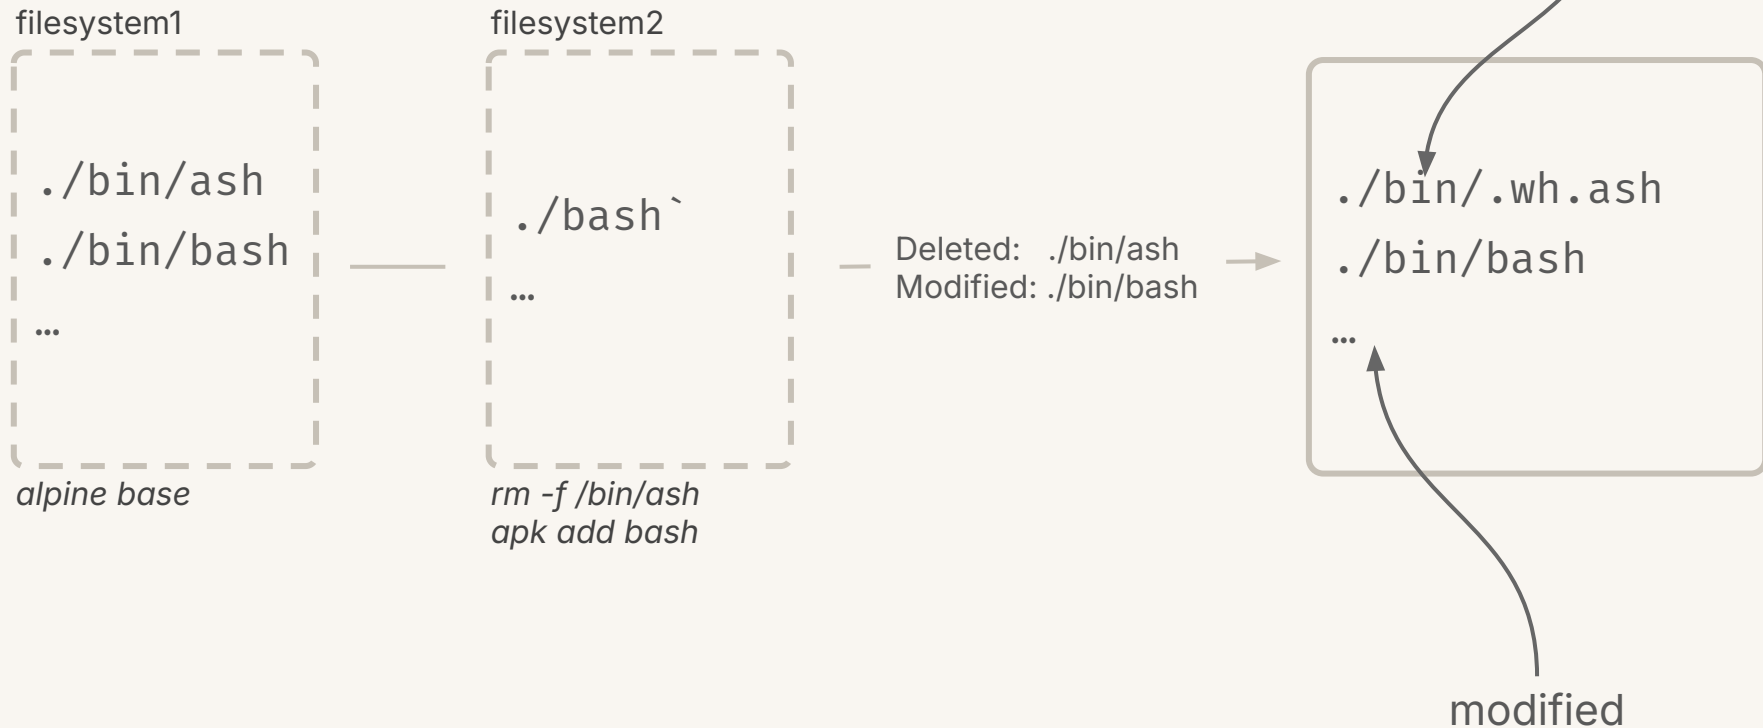
1. layer: create a changeset

```
FROM alpine
```

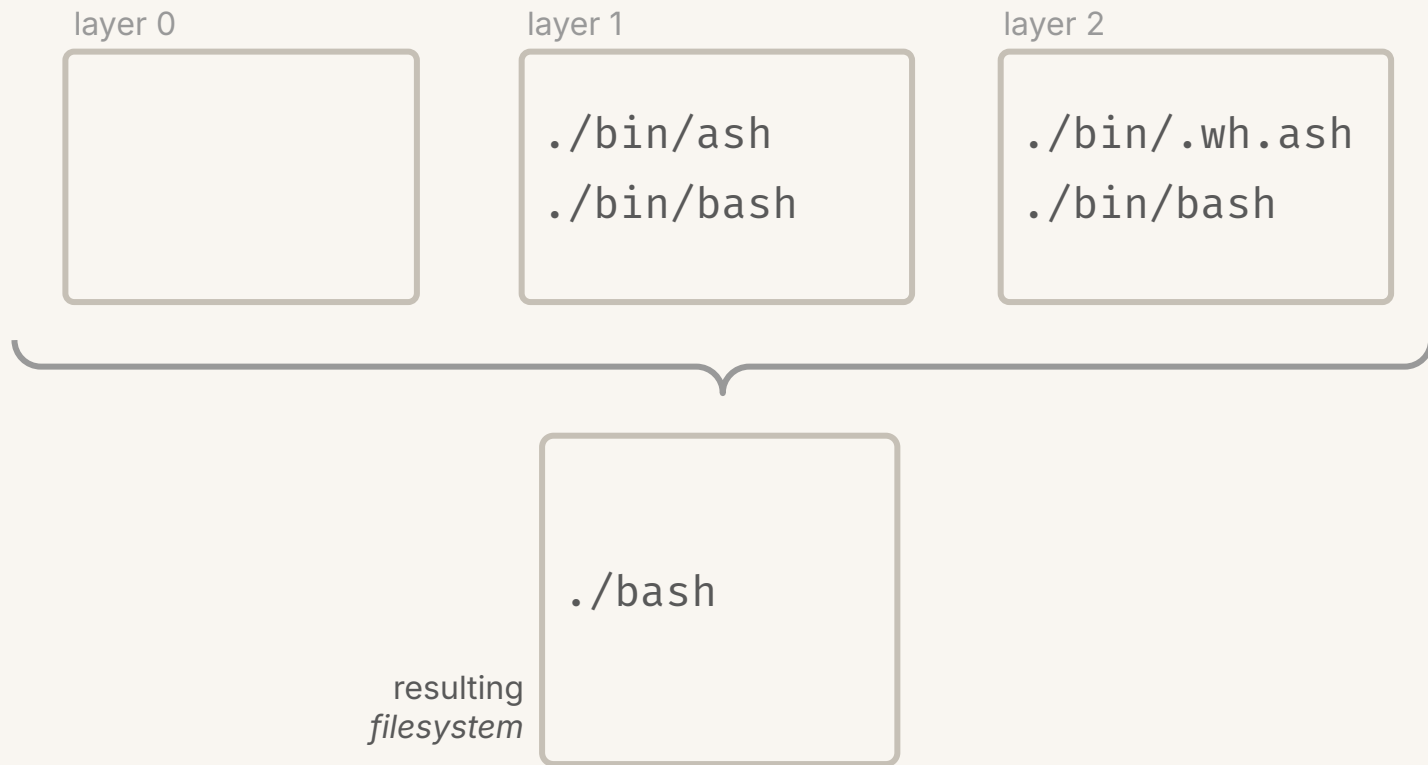
```
RUN rm -f /bin/ash \ # delete ash  
    && apk add bash  # update bash
```

```
...
```

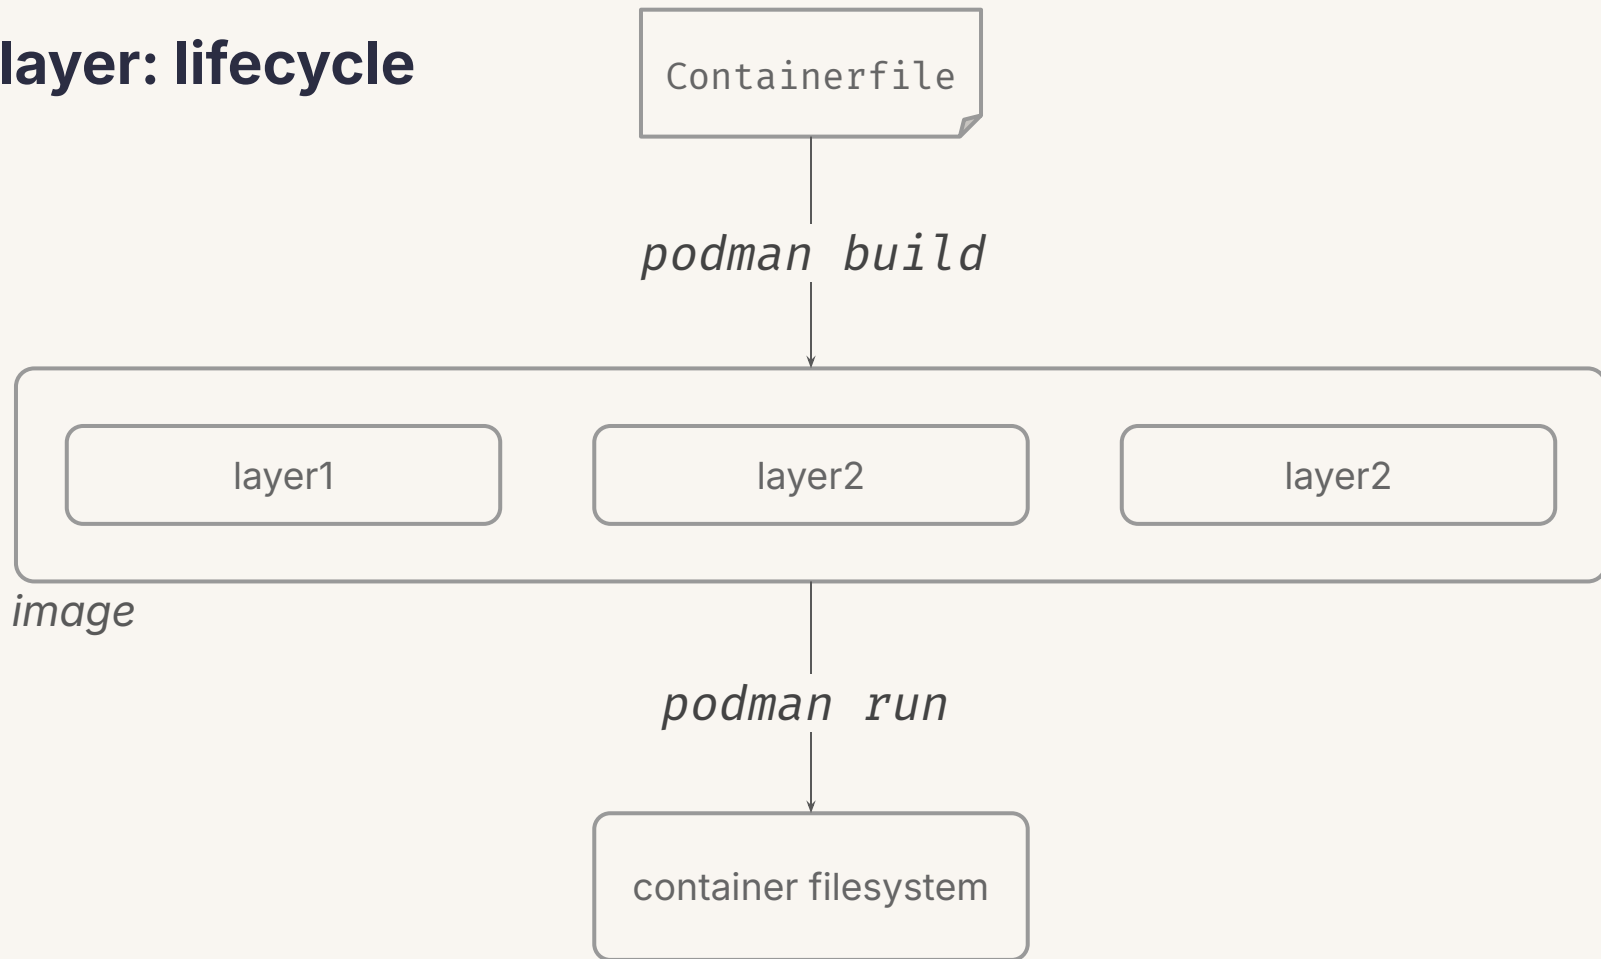
1. layer: create a changeset



1. layer: create the filesystem



1. layer: lifecycle



2. config

how to *run* the container

environment variables?


- HOSTNAME=globalhost
- PATH=/bin

container entrypoint

- ./start.sh

container cli options

- podman run \
 - volume
 - env
 - port



3. manifest


provide config *and* layer info

configuration digest


→ sha256:75b148a9a5b

layer digests

→ sha256:c37c06cdec9



content addressable store



content addressability

```
$ cat event.md  
Open Source Summit
```

```
$ sha256sum event.md
```

```
4c51d469b15d3423dd2d2...
```

```
$ mv event.md 4c51d469b15d3423dd2d2...
```


```
$ cat 4c51d469b15d3423dd2d2...
```

```
Open Source Summit
```

content address



content addressable store



4. index

manifest *for* manifest.json(s)

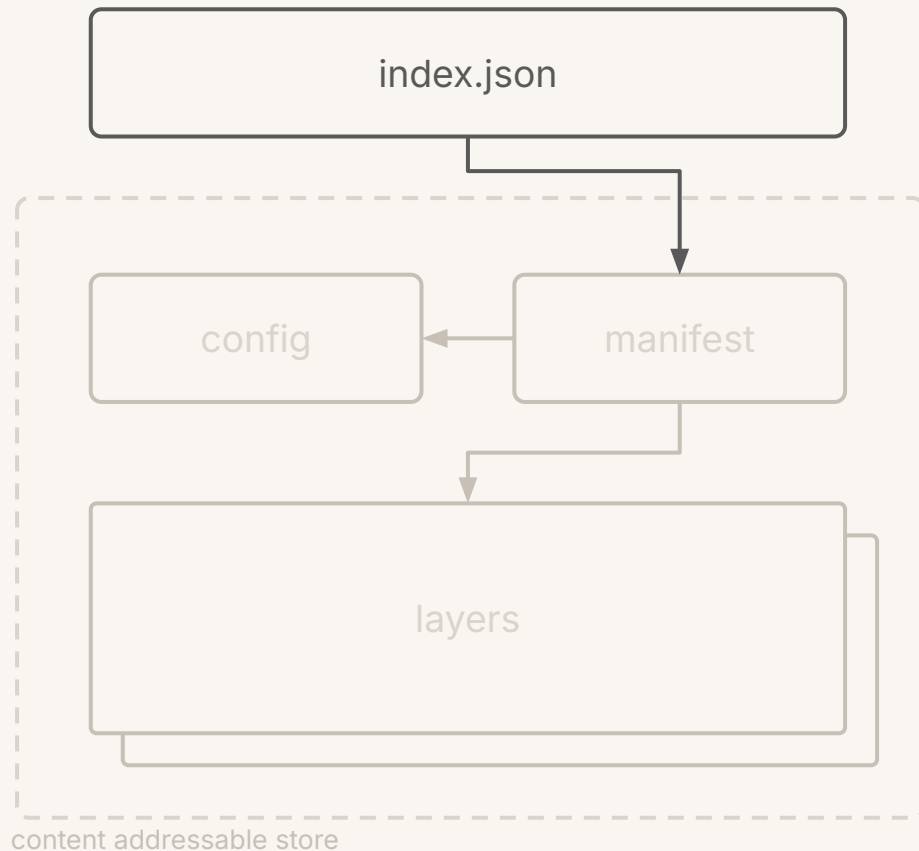
points to manifest.json

→ sha256:8289bd1bdc2


spans different platforms

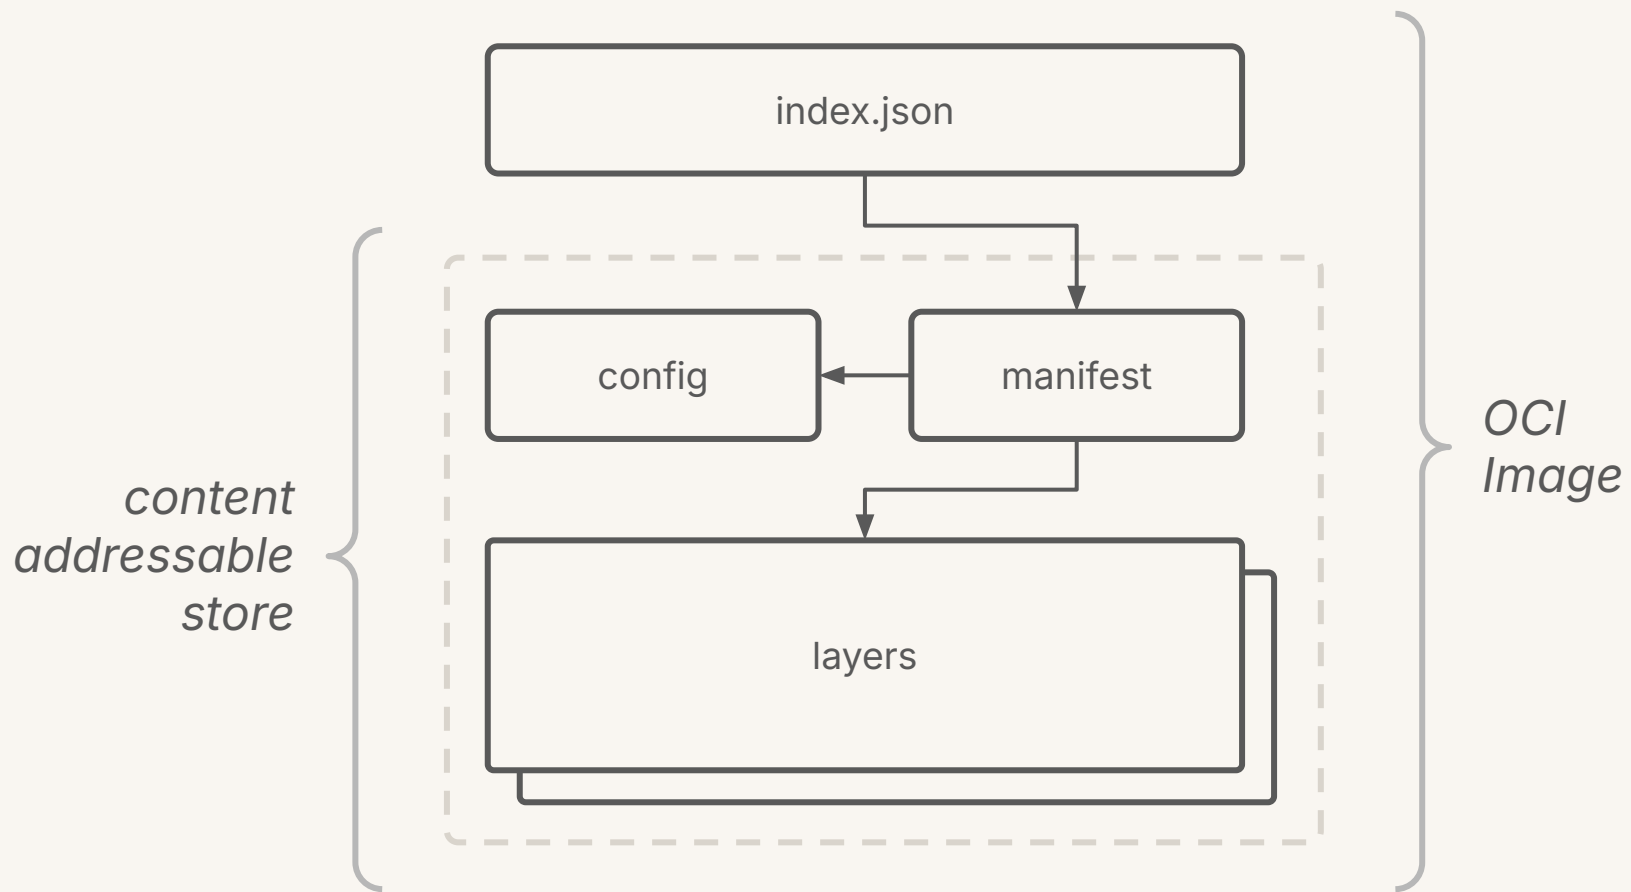
→ linux/amd64

→ linux/arm



multi-platform





let's build an image!

scratch image

first variant with *no* filesystem

no filesystem

FROM scratch

static binary

COPY ./hello ./

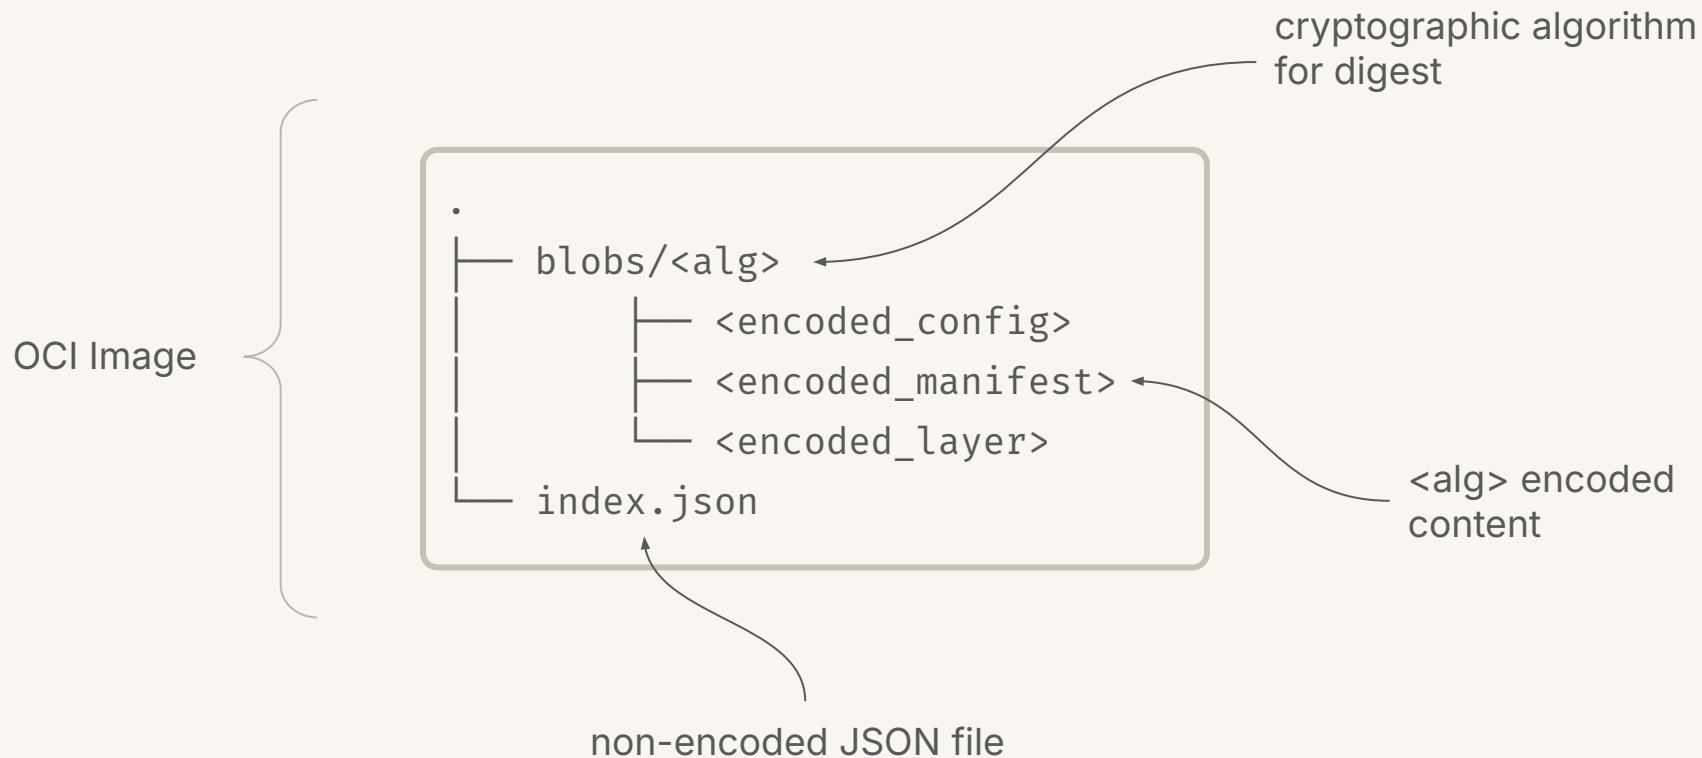
hello world
entrypoint

ENTRYPOINT ["./hello"]

podman build -f Containerfile .

image layout

the OCI image spec layout



scratch image: layer

first and only layer with a binary

gzipped
tar archive

```
$ gcc -o hello hello.c -static
$ tar -czvf layer.tar.gz hello

$ sha256sum layer.tar.gz
c37c06cdec9d6a0f2a2d5 layer.tar
$ mv layer.tar c37c06cdec9d6a0f2a2d5
```

content
addressable
archive

layout

```
$ tree ../..
├─ blobs/sha256
│   └─ c37c06cdec9d6a0f2a2d5
```

scratch image: config

setting the entrypoint and other config options

platform

entrypoint

```
$ cat config.json
{
  "architecture": "amd64",
  "os": "linux",
  "config": {
    "Entrypoint": [
      "./hello"
    ]
  }
}

$ mv config.json $(sha256sum
config.json | awk '{print $1}')
```

scratch image: manifest

identifying the layers and the config

```
$ cat manifest.json
```

```
{  
  "schemaVersion": 2,  
  "mediaType": "application/vnd.oci.image.manifest.v1+json",  
  "config": {  
    "mediaType": "application/vnd.oci.image.config.v1+json",  
    "digest": "sha256:bf7031d43f7d2c9ec77ed",  
    "size": 296  
  },  
  "layers": [  
    {  
      "mediaType": "application/vnd.oci.image.layer.v1.tar+gzip",  
      "digest": "sha256:583eb9106f1be6df fa3f0",  
      "size": 1377804  
    }  
  ]  
}
```

content
descriptor

config

layer digest

```
$ mv manifest.json $(sha256sum manifest.json | awk '{print $1}')
```


scratch image: index

identifying the only manifest in our image

manifest
digest

image
name:tag

```
{
  "schemaVersion": 2,
  "manifests": [
    {
      "mediaType": "application/vnd.oci.image.manifest.v1+json",
      "digest": "sha256:38f01cd4419e646946527",
      "size": 530,
      "annotations": {
        "org.opencontainers.image.ref.name": "hello:scratch"
      }
    }
  ]
}
```

scratch image: load & verify

```
$ tree
├── blobs/sha256
│   ├── 38f01cd4419e646946527
│   ├── c37c06cdec9d6a0f2a2d5
│   └── cd12bca58eb0ed29b24e4
└── index.json

$ tar -cf hello.tar *
```

• **\$ podman load < hello.tar**

```
Getting image source signatures
Copying blob c37c06cdec9d done    |
Copying config cd12bca58e done   |
Writing manifest to image destination
Loaded image: localhost/hello:scratch
```

\$ podman image ls hello

REPOSITORY	TAG	IMAGE ID	SIZE
hello	scratch	25e8b3bd9720	3.67MB

```
$ podman run localhost/hello:scratch world
hello world!
```

demo!

alpine base image

first variant but *with* alpine base image

minimal root
filesystem

static binary

`time`d
entrypoint

```
FROM alpine
```

```
COPY ./hello ./
```

```
ENTRYPOINT ["time", "./hello"]
```

```
podman build -f Containerfile .
```


alpine base: layer

using an alpine base root filesystem

```
$ wget https://.../alpine-minirootfs.tar.gz
$ sha256sum alpine-minirootfs.tar.gz
c59d5203bc6b8b6ef81f3
```

```
$ tree ../../..
├── blobs/sha256
│   ├── c37c06cdec9d6a0f2a2d5
│   ├── c59d5203bc6b8b6ef81f3
│   └── ...
```

content
addressable
archive



alpine base: manifest

Identifying both the layers and the config

```
"config": {  
  "mediaType": "application/vnd.oci.image.config.v1+json",  
  "digest": "sha256:bf8a32d43f7d2c9ec89fb",  
  "size": 296  
},  
"layers": [  
  {  
    "mediaType": "application/vnd.oci.image.layer.v1.tar+gzip",  
    "digest": "sha256:c59d5203bc6b8b6ef81f3",  
    "size": 3279768  
  },  
  {  
    "mediaType": "application/vnd.oci.image.layer.v1.tar+gzip",  
    "digest": "sha256:c37c06cdec9d6a0f2a2d5",  
    "size": 1372611  
  }  
]
```

layer hashes



alpine base: config

modify the entrypoint

entrypoint



```
{
  "architecture": "amd64",
  "os": "linux",
  "config": {
    "Entrypoint": [
      "time"
      "./hello"
    ]
  }
}
```

alpine base: index

updating the manifest digest and image tag

manifest
digest

image
name:tag

```
{
  "schemaVersion": 2,
  "manifests": [
    {
      "mediaType": "application/vnd.oci.image.manifest.v1+json",
      "digest": "sha256:38f01cd4419e646946527",
      "size": 530,
      "annotations": {
        "org.opencontainers.image.ref.name": "hello:alpine"
      }
    }
  ]
}
```


demo!

fin

→ container image *internals*

→ how *layers* are used

→ *multi-platform* images

resources

<https://github.com/opencontainers/image-spec>

<https://danishpraka.sh/posts/build-a-container-image-from-scratch/>

<https://danishpraka.sh/static/build-a-container-image-from-scratch.pdf>

danishpraka.sh

Software Engineer

SUSE